

# Off-Topic / Unsafe / Crisis Handling: Complete Guide

Revelia AI Astrologer Q&A · Build 27

## The core principle, stated once so it doesn't get lost

Classification and generation are two separate jobs, done by two separate model calls, and they should never be merged. Haiku's only job is to read a question and output one label. It never writes anything the user sees. The main answering model (Sonnet 5, Opus 4.8, or Fable 5) only ever runs on questions that are safe to answer. For crisis and clearly unsafe content, the main model is never called at all, the system responds directly with pre-written, hardcoded text. This is what makes the crisis path reliable: there's no point in the pipeline where a model is free-generating a helpline number.

## Step 1: The five categories, defined precisely

This is the part that was missing. A classifier can't work off a label name alone, it needs real definitions, or it'll guess inconsistently.

- **reflective** — Open personal questions about self, relationships, career, family, growth, patterns. Important calibration note: emotionally difficult topics like infidelity, breakups, divorce, grief, or conflict are reflective questions, not unsafe ones. "Is my partner cheating on me" or "should I confront my spouse about an affair" are completely normal relationship questions astrology apps answer every day. Don't let the topic's emotional weight get confused with actual unsafe content.
- **timing** — Binary, dated, or decision questions: "will X happen," "should I do Y," "when is a good time for Z."
- **off\_topic** — Unrelated to astrology or the user's own life (general trivia, coding help, requests with no connection to their chart or personal questions).
- **unsafe** — Sexually explicit content, anything sexual or romantic involving minors, requests for help with illegal activity (weapons, drugs, violence, other crime), or hate speech and harassment directed at a person or group. This is a narrow category. It is not "anything uncomfortable," it's specifically these things.
- **crisis** — Suicidal ideation or self-harm, however it's phrased, direct or indirect.

## Step 2: The actual classifier prompt (ready for Amey to use)

This is the missing piece. Drop this in as the system prompt for the Haiku 4.5 call:

*"You are a content classifier for an astrology Q&A app. Read the user's question and return exactly one label from this list: reflective, timing, off\_topic, unsafe, crisis. Do not answer the question. Do not explain your reasoning. Output only the label."*

*reflective: an open personal question about self, relationships, career, family, growth, or similar. This includes emotionally difficult topics like infidelity, breakups, grief, or conflict, these are still reflective questions.*

*timing: a binary, dated, or decision question, asking whether or when something will happen, or whether the user should take an action.*

*off\_topic: unrelated to astrology or the user's own life circumstances.*

*unsafe: sexually explicit content, anything sexual or romantic involving minors, requests for help with illegal activity, or hate speech and harassment directed at a person or group. A question about a difficult relationship or infidelity is reflective, not unsafe, unless it also contains one of the above.*

*crisis: expresses suicidal ideation or self-harm, regardless of phrasing, direct or indirect.*

*If uncertain between two labels, prefer the more cautious one: crisis over unsafe, unsafe over off\_topic, off\_topic over reflective."*

*That last line matters, it tells the model what to do when a question is genuinely ambiguous, default to the safer interpretation rather than guessing optimistically.*

## Step 3: Test fixtures, use these to validate the classifier before launch



If Amey's implementation misclassifies any of these, especially #3, #4 (infidelity ≠ unsafe) or #10 (indirect crisis language), that's a prompt calibration problem to fix before launch, not a fixture to change.

## Step 4: Routing logic, what happens after the label comes back

1. Question submitted by user
2. Gateway validation: reject blank/junk input (no model call, no credit, no cost)
3. Valid text → Haiku 4.5 classifier call
4. Classifier returns one label
5. IF label == crisis:
  - Do NOT call the main answering model at all
  - Return the hardcoded supportive response (Step 6)
  - Log the event (label + timestamp) for internal review
  - No credit deducted
6. IF label == unsafe:
  - Do NOT call the main answering model
  - Return the hardcoded generic decline (no explanation of why)
  - Log for abuse-pattern review
  - No credit deducted
7. IF label == off\_topic:
  - Do NOT call the main answering model
  - Return a hardcoded redirect message ("I'm built for questions about your own chart and life, try asking something like...")
  - No credit deducted
8. IF label == reflective OR timing:
  - Proceed to the normal pipeline (chart computation, prompt assembly, generation per tier)
  - Credit deducted on successful answer

This is a meaningful change from what's in the current docs, which left "sometimes reaches the model" undefined. Under this design, the main model is never invoked for crisis, unsafe, or off-topic questions, full stop. That's cheaper (one cheap Haiku call instead of two model calls) and safer (zero chance of the main model generating anything in a moment that matters).

## Step 5: Layer 2, the backstop, requires no extra building

If Haiku ever mislabels something, say, someone phrases a truly unsafe request in a way that slips through as "reflective", the main answering model still has its own built-in refusal behavior. It will decline

on its own without needing special instructions. The only build requirement here is a negative one: don't write the main system prompt in a way that overrides or suppresses that default behavior (for example, never instruct it to "always answer no matter what the question is"). Layer 2 is a safety net that exists for free as long as you don't engineer it away.

## Step 6: The crisis resource text (decided, no specific numbers)

**Final text:** "I know things feel really heavy right now. I'm not able to help with this one, but please don't go through it alone. Please reach out to a mental health crisis line in your area, they're trained for exactly this, or talk to someone you trust right now."

This was a deliberate choice, not a placeholder. Revelia's user base spans the US, India, Brazil, and Canada, among others, and a single hardcoded number (like the US's 988) would be wrong or unreachable for users elsewhere. Rather than building and maintaining a country-by-country number lookup, this response points to a local crisis line in general terms, which stays accurate everywhere without needing to track or verify per-country numbers.

**Worth a final confirmation from Sid:** that this general-purpose wording, rather than naming any specific number, is the intended final approach, not a stopgap. No further content is pending here unless that changes.

## Step 7: The off-topic and unsafe decline messages, lower stakes, safe for Amey to build directly

- **Unsafe:** "I can't help with that one, try a different question." No explanation, specificity just teaches people how to phrase around the block.
- **Off-topic:** something slightly warmer since it's not a safety issue, just a scope mismatch: "I'm built to help with questions about your own chart and life, not general topics. Try asking something like 'what does my chart say about...'"

## What this guide resolves versus what still needs a decision

**Resolved and ready for Amey to build directly:** the classifier prompt, the five category definitions, the routing logic, the test fixtures, the off-topic and unsafe decline messages, and the crisis response text.

**Still needs Sid's confirmation:** that skipping specific crisis helpline numbers in favor of general wording is the final intended approach, not a placeholder.